

Bin Ren

+39-329-324-8893 | bin.ren@unitn.it | <https://amazingren.github.io/>

 Amazingren |  Google Scholar |  LinkedIn

Trento, TN - 38123, Italy

RESEARCH INTERESTS

Representation Learning, 3D Scene Understanding, Multimodal Reasoning.

EDUCATION

• MBZUAI

Postdoc Associate

02, 2026 -

UAE

- Supervised by: Prof. Dr. Salman Khan and Prof. Dr. Fahad Shahbaz Khan .
- Topics: VLM(A), Multi-Modal Learning, Earth Observation

• University of Pisa (UNIPI) & University of Trento (UNITN)

National Ph.D. in AI for Society

11, 2021 - 10, 2025

Italy

- Supervised by: Prof. Dr. Nicu Sebe and Prof. Dr. Rita Cucchiara.
- Thesis Title: Representation Learning Via Transformer: From 2D to 3D

• Peking University (PKU)

MS in Computer Applied Technology

09, 2017 - 01. 2021

China

- Supervised by: Prof. Dr. Hong Liu.

- Thesis Title: Grouped Spatial-Temporal Enhancement Modeling Based Human Action Recognition
- Awards: The 2nd prize of the May 4th Youth Science Award (Challenge Cup of PKU) in 06, 2019

• Central South University (CSU)

BA in Microelectronics Manufacturing Engineering

09, 2012 - 07, 2016

China

- Scholarship: Three times a 3rd-class scholarship winner in 2013, 2014, and 2015.

VISITING RESEARCH EXPERIENCE

• Visiting Ph.D. Student at INSAIT [🌐]

Advised by Prof. Dr. Luc Van Gool

04,2024 - 09,2025

Sofia, Bulgaria

- *All-in-One Image Restoration:* At INSAIT, I led the development of efficient and scalable restoration models for diverse degradations. I proposed AnyIR (*Accepted* to TMLR), a lightweight framework leveraging shared semantics for unified restoration with >80% reduction in parameters and FLOPs. Building on this, I developed MIRAGE (*Accepted* to ICLR'26), which combines modular feature decomposition and manifold-based contrastive learning to enhance generalization. These works form a solid foundation for practical, high-quality image (and potentially video) enhancement under real-world conditions.
- *Large Multimodal Models (LMM):* I explored extending LMMs beyond natural images to structured, domain-specific modalities. We proposed EarthMind (*Submitted* to ICLR'26), a 4B vision-language framework for remote sensing, achieving SOTA performance on the proposed EarthMind-Bench and surpassing GPT-4o, demonstrating strong potential for generalization across multi-granular, multi-sensor visual domains.
- *3D Representation Learning:* Additionally, I am interested in processing irregular 3D data such as point clouds and Gaussian splats, with a focus on extracting meaningful representations to support downstream tasks. I contributed to ShapeSplat (**Accepted** by 3DV'25 as Oral), which introduces self-supervised learning on object-level 3D Gaussians via masked autoencoding. We then extended this to large-scale scenes with SceneSplat (**Accepted** to ICCV'25 as Oral), addressing structured scene-level representation learning. Then the dataset was extended from 7K to 49K scenes as SceneSplat++ (**Accepted** by NeurIPS'25).

• Visiting Ph.D. Student at Computer Vision Lab of ETH Zürich [🌐]

08, 2023 - 01, 2024

Zürich, Switzerland

Advised by Prof. Dr. Luc Van Gool

- *Low-Level Vision:* During my stay at the CVL of ETH Zürich, I proposed SemanIR, a transformer-based image restoration framework that constructs and shares key-semantic dictionaries to enable efficient and accurate restoration with linear attention complexity, achieving state-of-the-art performance across six benchmark tasks. This work has been **Accepted** by NeurIPS'24.
- *3D Representation Learning:* Meanwhile, I also worked on a novel framework that integrates contrastive learning into masked autoencoding for 3D point cloud pretraining by using dual random masking and shared-weight reconstruction, achieving decent performance across classification, segmentation, and few-shot tasks. This work has been **Accepted** by ACCV'24.

SELECTED PUBLICATIONS

C=CONFERENCE, J=JOURNAL, *=CORRESPONDING AUTHOR

- [C.1] Bin Ren, Runyi Yang, Qi Ma, Xu Zheng, Mengyuan Liu, Danda Pani Paudel, Luc Van Gool, Rita Cucchiara, Nicu Sebe. **Rethinking Expressivity and Degradation-Awareness in Attention for All-in-One Blind Image Restoration.** Accepted by ICLR'2026.
- [C.2] Bin Ren, Yawei Li, Xu Zheng, Yuqian Fu, Danda Pani Paudel, Hong Liu, Ming-Hsuan Yang, Luc Van Gool, and Nicu Sebe. **Manifold-aware Representation Learning for Degradation-agnostic Image Restoration.** Accepted by ICLR'2026.
- [C.3] Bin Ren, Xiaoshui Huang, Mengyuan Liu, Hong Liu, Fabio Poiesi, Nicu Sebe, and Guofeng Mei. **Masked Clustering Prediction for Unsupervised Point Cloud Pre-training.** Accepted by AAAI'2026 as **Oral**.
- [C.4] Bin Ren, Yawei Li, Jingyun Liang, Rakesh Ranjan, Mengyuan Liu, Rita Cucchiara, Luc Van Gool, Ming-Hsuan Yang, and Nicu Sebe. **Sharing Key Semantics in Transformer Makes Efficient Image Restoration.** Accepted by NeurIPS'2024.
- [C.5] Bin Ren, Yahui Liu, Yue Song, Wei Bi, Rita Cucchiara, Nicu Sebe, and Wei Wang. **Masked Jigsaw Puzzle: A Versatile Position Embedding for Vision Transformers.** Accepted by CVPR'2023.
- [C.6] Bin Ren, Guofeng Mei, Danda Pani Paudel, Weijie Wang, Yawei Li, Mengyuan Liu, Rita Cucchiara, Luc Van Gool, and Nicu Sebe. **Bringing masked autoencoders explicit contrastive properties for point cloud self-supervised learning.** Accepted by ACCV'2024.
- [C.7] Bin Ren, Hao Tang, Yiming Wang, Xia Li, Wei Wang, and Nicu Sebe. **PI-Trans: Parallel-ConvMLP and Implicit-Transformation Based GAN for Cross-View Image Translation.** Accepted by ICASSP'2023.
- [C.8] Bin Ren, Hao Tang, and Nicu Sebe. **Cascaded Cross MLP-Mixer GANs for Cross-View Image Translation.** Accepted by BMVC'2021 as **Oral**.
- [C.9] Davide Lobba, Fulvio Sanguigni, Bin Ren*, Marcella Cornia, Rita Cucchiara, and Nicu Sebe. **Inverse Virtual Try-On: Generating Multi-Category Product-Style Images from Clothed Individuals.** Accepted by ICLR'2026.
- [C.10] Mengjiao Ma, Qi Ma, Yue Li, Jiahuan Cheng, Runyi Yang, Bin Ren*, Nikola Popovic, Mingqiang Wei, Nicu Sebe, Luc Van Gool, Theo Gevers, Martin R Oswald, Danda Pani Paudel. **SceneSplat++: A Large Dataset and Comprehensive Benchmark for Language Gaussian Splatting.** Accepted by NeurIPS'2025.
- [C.11] Yue Li, Qi Ma, Runyi Yang, Huapeng Li, Mengjiao Ma, Bin Ren*, Nikola Popovic, Nicu Sebe, Ender Konukoglu, Theo Gevers, Luc Van Gool, Martin R Oswald, Danda Pani Paudel. **Scenesplat: Gaussian splatting-based scene understanding with vision-language pretraining.** Accepted by ICCV'2025 as **Oral**.
- [C.12] Yuqian Fu, Runze Wang, Bin Ren, Guolei Sun, Biao Gong, Yanwei Fu, Danda Pani Paudel, Xuanjing Huang, Luc Van Gool **ObjectRelator: Enabling Cross-View Object Relation Understanding Across Ego-Centric and Exo-Centric Perspectives.** Accepted by ICCV'2025 as **Highlight**.
- [C.13] Qi Ma, Yue Li, Bin Ren*, Nicu Sebe, Ender Konukoglu, Theo Gevers, Luc Van Gool, and Danda Pani Paudel. **ShapeSplat: A Large-scale Dataset of Gaussian Splats and Their Self-Supervised Pretraining.** Accepted by 3DV'2025 as **Oral**.
- [C.14] Mengyi Zhao, Mengyuan Liu, Bin Ren, Shuling Dai, and Nicu Sebe. **Modiff: Action-conditioned 3d motion generation with denoising diffusion probabilistic models.** Accepted by ICASSP'2024.
- [C.15] Chang Liu, Mengyi Zhao, Bin Ren, Mengyuan Liu, and Nicu Sebe. **Spatio-Temporal Graph Diffusion for Text-Driven Human Motion Generation.** Accepted by BMVC'2023 as **Oral**.
- [J.1] Bin Ren, Hao Tang, Fanyang Meng, Runwei Ding, Philip Torr, and Nicu Sebe. **Cloth interactive transformer for virtual try-on.** Accepted by ACM ToMM, 2023.
- [J.2] Weixin Ye, Wei Wang, Yahui Liu, Yue Song, Bin Ren*, Wei Bi, Rita Cucchiara, and Nicu Sebe. **A Unified Masked Jigsaw Puzzle Framework for Vision and Language Models.** Accepted by TPAMI, 2025.
- [J.3] Hao Tang, Bin Ren, and Nicu Sebe. **A pure MLP-Mixer-based GAN framework for guided image translation.** Accepted by PR, 2024.
- [J.4] Ti Wang, Mengyuan Liu, Hong Liu, Bin Ren, Yingxuan You, Wenhao Li, Nicu Sebe, and Xia Li. **Uncertainty-Aware Testing-Time Optimization for 3D Human Pose Estimation.** Accepted by TMM, 2025.
- [J.5] Wenhao Li, Mengyuan Liu, Hong Liu, Bin Ren, Xia Li, Yingxuan You, and Nicu Sebe. **HYRE: Hybrid Regressor for 3D Human Pose and Shape Estimation.** Accepted by TIP, 2024.

SELECTED SUBMITTED WORKS

C=CONFERENCE, J=JOURNAL, *=CORRESPONDING AUTHOR

- [J.1] Bin Ren, Eduard Zamfir, Zongwei Wu, Yawei Li, Yidi Li, Danda Pani Paudel, Radu Timofte, Ming-Hsuan Yang, Luc Van Gool, and Nicu Sebe. **Any Image Restoration via Efficient Spatial-Frequency Degradation Adaptation.** Accepted with Minor Revision from TMLR.
- [C.2] Yan Shu, Bin Ren*, Zhitong Xiong, Danda Pani Paudel, Luc Van Gool, Begum Demir, Nicu Sebe, Paolo Rota. **EarthMind: Towards Multi-Granular and Multi-Sensor Earth Observation with Large Multimodal Models.** Submitted to ICLR in 2025.

ACADEMIC SERVICES

- Workshop/Challenge Organizer: NTIRE (9th, 10th) at CVPR'24,25), AIM (6th) at ICCV'25, SPI (1st) at SWC'24.
- Conference Reviewer: NeurIPS, ICML, ICLR, CVPR, ICCV, ECCV, ACCV, BMVC, 3DV, and ICASSP, etc.
- Journal Reviewer: TPAMI, TIP, IJCV, RAL, TMM, and ToMM, etc.
- Area Chair: BMVC'25

SKILLS

- Programming: Python, PyTorch, Linux, Huggingface, Git, ViM.
- Research: Vision Transformer, Low-Level Vision, Gaussian Splatting, 3D Vision, Self-Supervised Learning.

ADDITIONAL INFORMATION

Languages: English (C1, Frequent use), Italian (A2, Basic), Mandarin (Native).

Interests: Cooking, Hiking, and Singing.